

# Question #02: Linear Regression

## 1. Data Preprocessing and Exploration

### Load the Dataset and Handle Missing Values

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler

df = pd.read_csv("linear_regression_dataset.csv")

print(df.info())
print(df.isnull().sum())

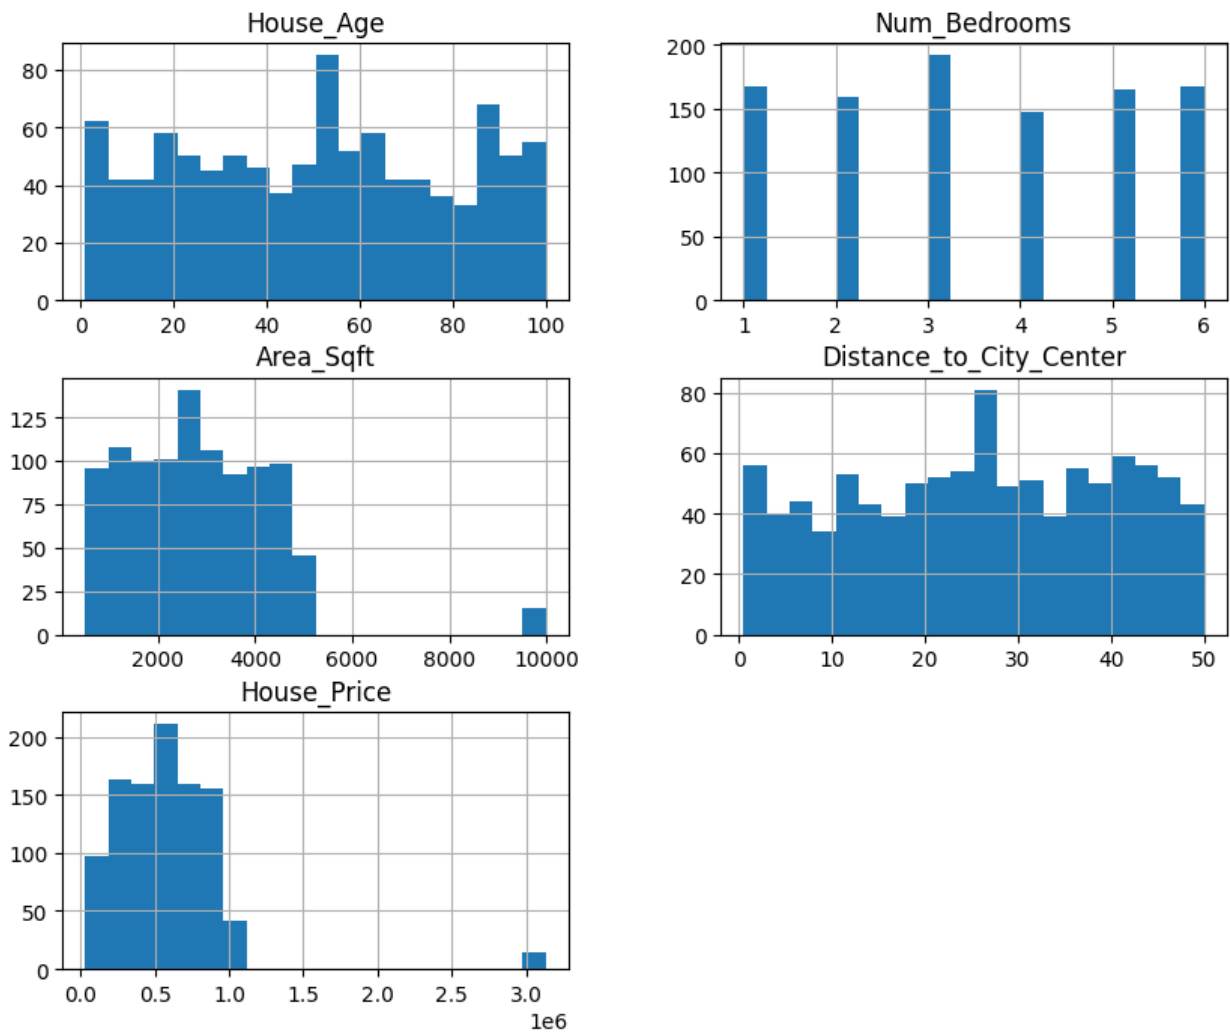
df.fillna(df.median(numeric_only=True), inplace=True)

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 5 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   House_Age                             968 non-null    float64
1   Num_Bedrooms                           973 non-null    float64
2   Area_Sqft                             965 non-null    float64
3   Distance_to_City_Center                981 non-null    float64
4   House_Price                            968 non-null    float64
dtypes: float64(5)
memory usage: 39.2 KB
None
House_Age                32
Num_Bedrooms              27
Area_Sqft                 35
Distance_to_City_Center   19
House_Price               32
dtype: int64
```

### Visualizing Feature Distributions


```
df.hist(figsize=(10, 8), bins=20)
plt.suptitle("Feature Distributions")
plt.show()
```

## Feature Distributions



## Feature Correlation Analysis

```
plt.figure(figsize=(8, 6))
sns.heatmap(df.corr(), annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Feature Correlation Heatmap")
plt.show()
```



## 2. Feature Engineering and Selection

### Feature Normalization & Correlation Analysis

```
from sklearn.preprocessing import PolynomialFeatures

scaler = StandardScaler()
X_scaled =
pd.DataFrame(scaler.fit_transform(df.drop(columns=['House_Price'])),
columns=df.columns[:-1])

correlations = df.corr()['House_Price'].drop('House_Price')
print("Feature Correlations with House_Price:\n", correlations)
```

```
selected_features = correlations[abs(correlations) >
0.2].index.tolist()
X_selected = X_scaled[selected_features]
```

```
Feature Correlations with House_Price:
House_Age          0.005886
Num_Bedrooms       0.032417
Area_Sqft          0.939058
Distance_to_City_Center -0.031151
Name: House_Price, dtype: float64
```

## Create Polynomial Features

```
poly = PolynomialFeatures(degree=2, include_bias=False)
X_poly = poly.fit_transform(X_selected)

print("Polynomial Features Shape:", X_poly.shape)

Polynomial Features Shape: (1000, 2)
```

## 3. Train a Linear Regression Model

### Splitting Data into Training & Testing Sets

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score

X = X_poly
y = df['House_Price']

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

### Training & Evaluating Linear Regression Model

```
model = LinearRegression()
model.fit(X_train, y_train)

y_pred = model.predict(X_test)
```

```

mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Linear Regression Performance:\nMAE: {mae}\nMSE: {mse}\nR2 Score: {r2}")

Linear Regression Performance:
MAE: 47552.10709897928
MSE: 7035494815.528621
R2 Score: 0.9723984550181082

```

## 4. Implement Linear Regression using Gradient Descent

### Gradient Descent Implementation

```

import numpy as np

m, n = X_train.shape
theta = np.zeros(n)
alpha = 0.01
iterations = 1000

def gradient_descent(X, y, theta, alpha, iterations):
    m = len(y)
    for i in range(iterations):
        predictions = X.dot(theta)
        errors = predictions - y
        gradient = (1/m) * X.T.dot(errors)
        theta -= alpha * gradient
    return theta

```

### Training Gradient Descent Model

```

# Add bias column (intercept) to X
X_train_bias = np.c_[np.ones((X_train.shape[0], 1)), X_train]
X_test_bias = np.c_[np.ones((X_test.shape[0], 1)), X_test]

y_train_np = y_train.values

theta_final = gradient_descent(X_train_bias, y_train_np,
                                np.zeros(X_train_bias.shape[1]), alpha, iterations)

```

```

y_pred_gd = X_test_bias.dot(theta_final)

mae_gd = mean_absolute_error(y_test, y_pred_gd)
mse_gd = mean_squared_error(y_test, y_pred_gd)
r2_gd = r2_score(y_test, y_pred_gd)

print(f"Gradient Descent Regression Performance:\nMAE: {mae_gd}\nMSE: {mse_gd}\nR2 Score: {r2_gd}")

```

Gradient Descent Regression Performance:  
MAE: 47602.07332171253  
MSE: 7030830459.975481  
R<sup>2</sup> Score: 0.9724167541460285

## 5. Predict House Prices for New Data

### Scaling and Transformation for New Data

```

import numpy as np
import pandas as pd
from sklearn.preprocessing import StandardScaler, PolynomialFeatures

scaler_single = StandardScaler()
scaler_single.fit(df[['Area_Sqft']])

poly_single = PolynomialFeatures(degree=2, include_bias=False)
poly_single.fit(df[['Area_Sqft']])

def predict_house_price(new_data):
    if not isinstance(new_data, list):
        raise ValueError("Input must be a list containing one feature value.")

    # Convert new_data into a DataFrame with correct column name
    new_data_df = pd.DataFrame([new_data], columns=['Area_Sqft'])

    new_data_scaled = scaler_single.transform(new_data_df)

    new_data_poly = poly_single.transform(new_data_scaled)

    predicted_price = model.predict(new_data_poly)[0]

    return predicted_price

```

# Predicting House Prices

```
print("Feature used during training:", ['Area_Sqft'])
```

```
new_house_1 = [2000]
new_house_2 = [1500]
new_house_3 = [3000]
```

```
predicted_price_1 = predict_house_price(new_house_1)
predicted_price_2 = predict_house_price(new_house_2)
predicted_price_3 = predict_house_price(new_house_3)
```

```
print(f"Predicted House Price for 2000 sqft: ${predicted_price_1:.2f}")
print(f"Predicted House Price for 1500 sqft: ${predicted_price_2:.2f}")
print(f"Predicted House Price for 3000 sqft: ${predicted_price_3:.2f}")
```

```
Feature used during training: ['Area_Sqft']
Predicted House Price for 2000 sqft: $380809.10
Predicted House Price for 1500 sqft: $302954.74
Predicted House Price for 3000 sqft: $568047.87
```

```
C:\Users\HP_840G4\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\base.py:493: UserWarning: X does not have valid feature names, but PolynomialFeatures was fitted with feature names
  warnings.warn(
```

```
C:\Users\HP_840G4\AppData\Local\Programs\Python\Python312\Lib\site-  
packages\sklearn\base.py:493: UserWarning: X does not have valid  
feature names, but PolynomialFeatures was fitted with feature names  
warnings.warn(
```

```
C:\Users\HP_840G4\AppData\Local\Programs\Python\Python312\Lib\site-  
packages\sklearn\base.py:493: UserWarning: X does not have valid  
feature names, but PolynomialFeatures was fitted with feature names  
warnings.warn(
```